Langfuse Output Summary

Paul Terry, Hayden Hanes, Will Hansen

Tracing Images Below, along with an explanation below.

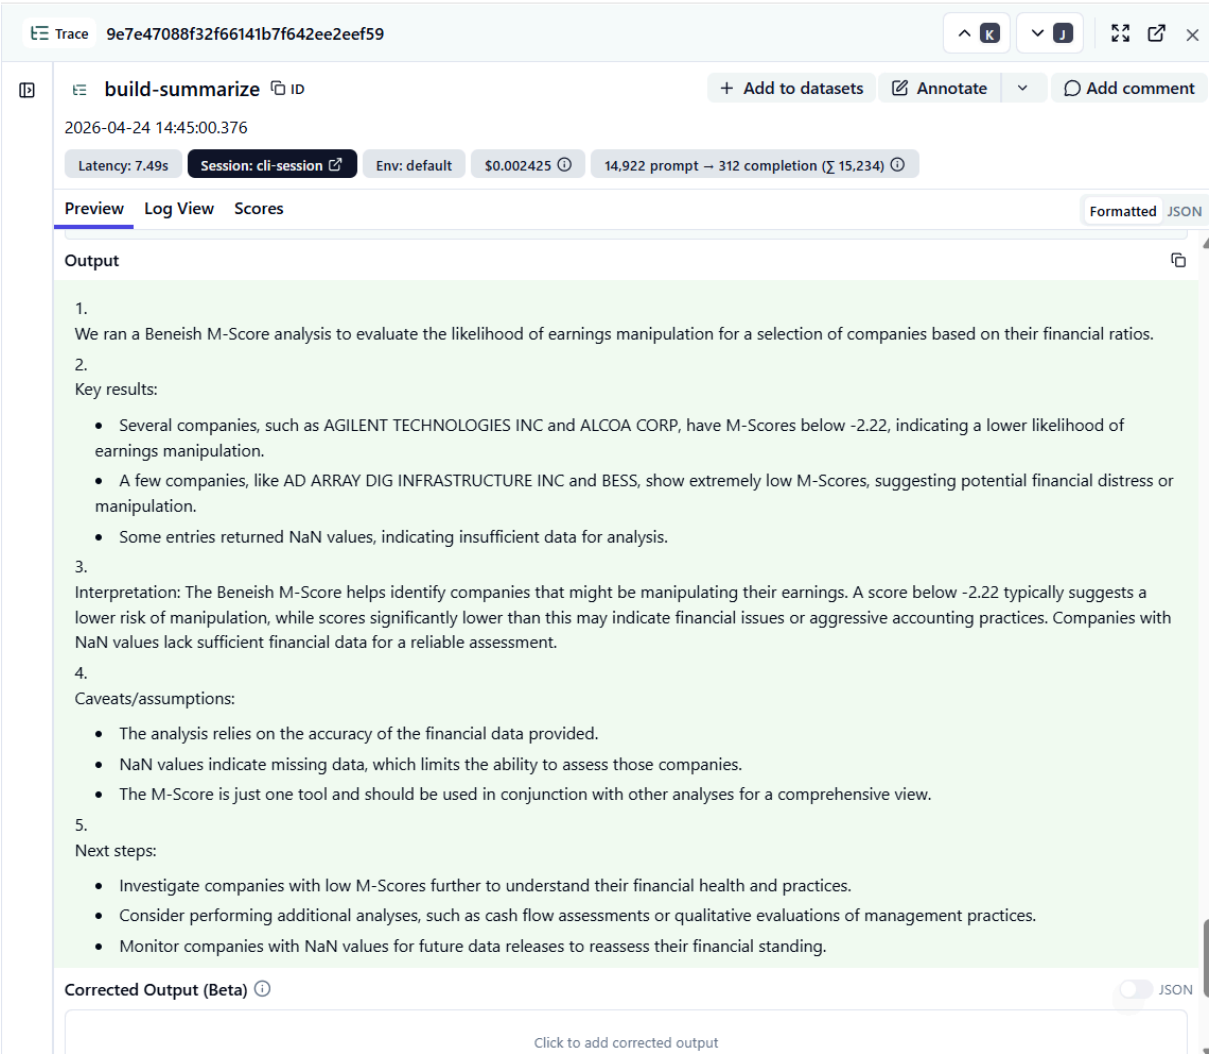

Trace a39b39b8b659af8a20e22e854c28dbc0

build-summarize ID

+ Add to datasets | Annotate | Add comment

2026-04-24 14:45:27.273

Latency: 9.42s | Session: cli-session | Env: default | $0.002502 | 15,437 prompt → 311 completion (∑ 15,748)

Preview | Log View | Scores

Formatted JSON

Output

1.
We ran a distress scoring analysis on various companies, providing a score that indicates their financial distress level based on the provided data.

2.
Key results:

- The scores range from extremely low values (e.g., 9.496464e-12 for Agilent Technologies Inc) to higher values (e.g., 0.999987 for Adial Pharmaceuticals).
- Many companies have scores close to zero, indicating low distress, while a few have scores approaching one, suggesting high distress.

3.
Interpretation: The distress scores indicate the financial health of the companies analyzed. A score close to zero suggests that a company is in good financial standing, while a score closer to one indicates significant financial distress. For example, Agilent Technologies has a very low score, suggesting it is financially stable, whereas Adial Pharmaceuticals has a high score, indicating potential financial trouble.

4.
Caveats/assumptions:

- The scores are based on a specific model and may not capture all aspects of a company's financial health.
- Some companies have missing data (NaN), which could affect the overall analysis and interpretation.
- The context of the scores (e.g., industry conditions, market trends) is not considered in this analysis.

5.
Next steps:

- Investigate the companies with high distress scores to understand the underlying reasons for their financial situation.
- Consider conducting a more detailed financial analysis, including reviewing financial statements and market conditions.
- Monitor the distress scores over time to identify trends and changes in financial health.

Corrected Output (Beta) JSON

Trace beb01c0b612c1a5b9b10feaa42da4b30

build-summarize ID

+ Add to datasets | Annotate | Add comment

2026-04-24 14:49:49.642

Latency: 9.53s | Session: cli-session | Env: default | $0.002366 | 14,943 prompt → 208 completion (∑ 15,151)

Preview | Log View | Scores

Formatted JSON

Output

1.
We analyzed the likelihood of earnings manipulation for Apple Inc. (AAPL) using a probability model that assesses financial metrics.

2.
Key results:

- Probability of manipulation (PROBM): -2.404749
- Manipulation metric (PMAN): 8.091787e-03

3.
Interpretation: The negative probability score indicates that Apple Inc. is less likely to engage in earnings manipulation compared to companies with higher positive scores. The low manipulation metric suggests that the risk of financial misrepresentation is relatively low for Apple.

4.
Caveats/assumptions:

- The model relies on historical data and may not account for recent changes in company practices or external factors.
- The results are based on statistical probabilities and do not guarantee future behavior.

5.
Next steps:

- Monitor Apple's financial reports and disclosures for any signs of unusual accounting practices.
- Conduct a deeper analysis of specific financial metrics that could indicate manipulation.
- Compare Apple's results with industry peers to assess relative risk levels.

Corrected Output (Beta) JSON

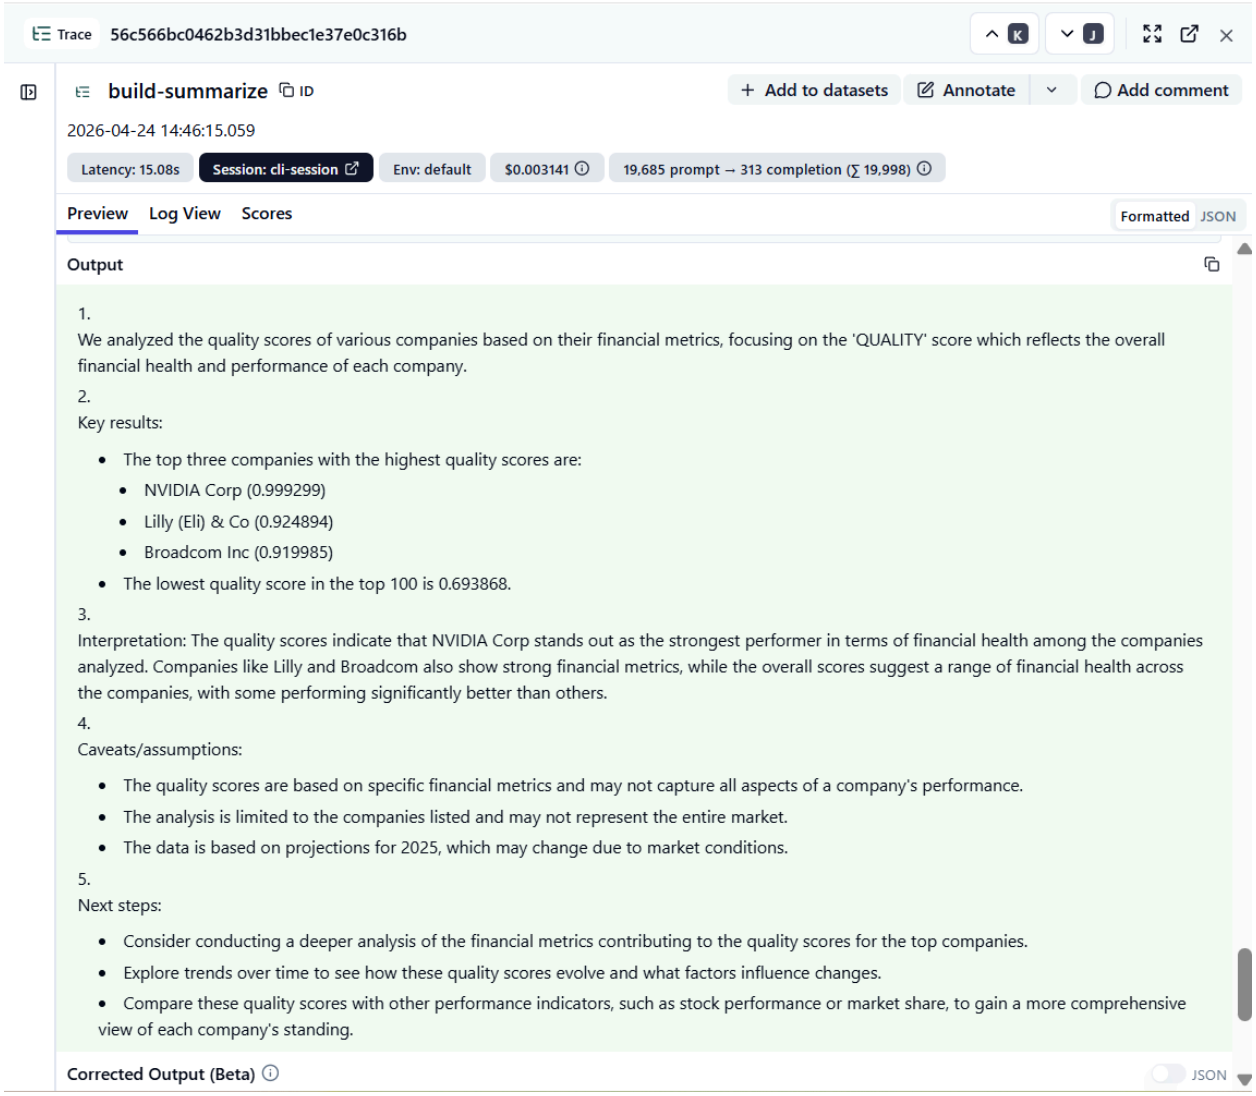

Tracing and Error Log Summary

The Langfuse traces show how our agent processed multiple financial analysis tasks, including earnings manipulation detection (Beneish M-score), distress scoring, probability modeling, and quality scoring. Overall, our agent performed very well. Most of our output was detailed, qualitative financial information. We hope to add more quantitative financial analysis to our agent, and we get those suggestions from the output itself. Some outputs returned NaN values, indicating missing or incomplete

financial data. Certain analyses relied heavily on qualitative interpretation, even when numerical outputs were available.

Output Explanation

The agent generated several types of financial analysis:

- Beneish M-Score Analysis:

  Used to detect potential earnings manipulation. Scores below -2.22 indicate a lower likelihood of manipulation, while higher scores suggest potential risk. The agent identified companies with both low-risk and potentially concerning scores
- Distress Scoring Model:

  Produced values between 0 and 1, where values closer to 1 indicate higher financial distress. Most companies showed low distress, but a few had elevated risk levels.
- Probability Model

  The agent calculated a negative manipulation probability score for Apple, suggesting a low likelihood of earnings manipulation.
- Quality Score Analysis:

  Ranked companies based on overall financial strength. NVIDIA, Eli Lilly, and Broadcom were identified as top performers, indicating strong financial health

RAG Performance Evaluation

In our agent, RAG had somewhat of a minimal impact on agent performance, and in some cases, slightly degraded it. The primary reason is that the agent's tasks were numerical and model-based rather than knowledge-retrieval-based. The outputs were mostly generated from structured financial data rather than

external documents. As a result, RAG did not have much of an impact on enhancing the agent's ability to perform calculations or improve accuracy.

Improvements

RAG performance could be improved by integrating more relevant and structured data sources, particularly those that go beyond the basic financial ratios currently used in the analysis. These metrics would allow the agent to move beyond surface-level analysis and produce more robust, quantitative insights. We could also improve RAG by incorporating historical and time-series data. Instead of analyzing companies at a single point in time, the agent could retrieve, multi-year financial trends and historical performance during economic downturns. One other idea we could integrate would be more anecdotal and qualitative evidence, which could help enhance interpretation.